

Lab 9: Mini Project

Vivian Cai

2/12/2022

1. Exploratory Data Analysis

Before starting the project, I downloaded the WisconsinCancer.csv file from the class website and moved it into the Lab9 folder (same directory as my R markdown file)

```
# Save input data file into Project directory
fna.data <- "WisconsinCancer.csv"

# input the data and store as wisc.df
wisc.df <- read.csv(fna.data, row.names=1)
head(wisc.df)
```

```
##      diagnosis radius_mean texture_mean perimeter_mean area_mean
## 842302         M      17.99      10.38      122.80      1001.0
## 842517         M      20.57      17.77      132.90      1326.0
## 84300903        M      19.69      21.25      130.00      1203.0
## 84348301         M      11.42      20.38       77.58       386.1
## 84358402         M      20.29      14.34      135.10      1297.0
## 843786          M      12.45      15.70       82.57       477.1
##      smoothness_mean compactness_mean concavity_mean concave.points_mean
## 842302      0.11840      0.27760      0.3001      0.14710
## 842517      0.08474      0.07864      0.0869      0.07017
## 84300903      0.10960      0.15990      0.1974      0.12790
## 84348301      0.14250      0.28390      0.2414      0.10520
## 84358402      0.10030      0.13280      0.1980      0.10430
## 843786      0.12780      0.17000      0.1578      0.08089
##      symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
## 842302      0.2419      0.07871      1.0950      0.9053      8.589
## 842517      0.1812      0.05667      0.5435      0.7339      3.398
## 84300903      0.2069      0.05999      0.7456      0.7869      4.585
## 84348301      0.2597      0.09744      0.4956      1.1560      3.445
## 84358402      0.1809      0.05883      0.7572      0.7813      5.438
## 843786      0.2087      0.07613      0.3345      0.8902      2.217
##      area_se smoothness_se compactness_se concavity_se concave.points_se
## 842302     153.40      0.006399      0.04904      0.05373      0.01587
## 842517      74.08      0.005225      0.01308      0.01860      0.01340
## 84300903     94.03      0.006150      0.04006      0.03832      0.02058
## 84348301     27.23      0.009110      0.07458      0.05661      0.01867
## 84358402     94.44      0.011490      0.02461      0.05688      0.01885
## 843786     27.19      0.007510      0.03345      0.03672      0.01137
##      symmetry_se fractal_dimension_se radius_worst texture_worst
```

```
## 842302      0.03003      0.006193      25.38      17.33
## 842517      0.01389      0.003532      24.99      23.41
## 84300903    0.02250      0.004571      23.57      25.53
## 84348301    0.05963      0.009208      14.91      26.50
## 84358402    0.01756      0.005115      22.54      16.67
## 843786      0.02165      0.005082      15.47      23.75
##           perimeter_worst area_worst smoothness_worst compactness_worst
## 842302      184.60      2019.0      0.1622      0.6656
## 842517      158.80      1956.0      0.1238      0.1866
## 84300903    152.50      1709.0      0.1444      0.4245
## 84348301      98.87      567.7      0.2098      0.8663
## 84358402    152.20      1575.0      0.1374      0.2050
## 843786      103.40      741.6      0.1791      0.5249
##           concavity_worst concave.points_worst symmetry_worst
## 842302      0.7119      0.2654      0.4601
## 842517      0.2416      0.1860      0.2750
## 84300903    0.4504      0.2430      0.3613
## 84348301    0.6869      0.2575      0.6638
## 84358402    0.4000      0.1625      0.2364
## 843786      0.5355      0.1741      0.3985
##           fractal_dimension_worst
## 842302      0.11890
## 842517      0.08902
## 84300903    0.08758
## 84348301    0.17300
## 84358402    0.07678
## 843786      0.12440
```

```
dim(wisc.df)
```

```
## [1] 569 31
```

```
# Creating a new data frame what omits the diagnosis column
wisc.data <- wisc.df[,-1]
head(wisc.data)
```

```
##           radius_mean texture_mean perimeter_mean area_mean smoothness_mean
## 842302      17.99      10.38      122.80      1001.0      0.11840
## 842517      20.57      17.77      132.90      1326.0      0.08474
## 84300903    19.69      21.25      130.00      1203.0      0.10960
## 84348301    11.42      20.38      77.58      386.1      0.14250
## 84358402    20.29      14.34      135.10      1297.0      0.10030
## 843786      12.45      15.70      82.57      477.1      0.12780
##           compactness_mean concavity_mean concave.points_mean symmetry_mean
## 842302      0.27760      0.3001      0.14710      0.2419
## 842517      0.07864      0.0869      0.07017      0.1812
## 84300903    0.15990      0.1974      0.12790      0.2069
## 84348301    0.28390      0.2414      0.10520      0.2597
## 84358402    0.13280      0.1980      0.10430      0.1809
## 843786      0.17000      0.1578      0.08089      0.2087
##           fractal_dimension_mean radius_se texture_se perimeter_se area_se
## 842302      0.07871      1.0950      0.9053      8.589 153.40
## 842517      0.05667      0.5435      0.7339      3.398 74.08
```

```
## 84300903      0.05999      0.7456      0.7869      4.585      94.03
## 84348301      0.09744      0.4956      1.1560      3.445      27.23
## 84358402      0.05883      0.7572      0.7813      5.438      94.44
## 843786        0.07613      0.3345      0.8902      2.217      27.19
##      smoothness_se compactness_se concavity_se concave.points_se
## 842302      0.006399      0.04904      0.05373      0.01587
## 842517      0.005225      0.01308      0.01860      0.01340
## 84300903      0.006150      0.04006      0.03832      0.02058
## 84348301      0.009110      0.07458      0.05661      0.01867
## 84358402      0.011490      0.02461      0.05688      0.01885
## 843786        0.007510      0.03345      0.03672      0.01137
##      symmetry_se fractal_dimension_se radius_worst texture_worst
## 842302      0.03003      0.006193      25.38      17.33
## 842517      0.01389      0.003532      24.99      23.41
## 84300903      0.02250      0.004571      23.57      25.53
## 84348301      0.05963      0.009208      14.91      26.50
## 84358402      0.01756      0.005115      22.54      16.67
## 843786        0.02165      0.005082      15.47      23.75
##      perimeter_worst area_worst smoothness_worst compactness_worst
## 842302      184.60      2019.0      0.1622      0.6656
## 842517      158.80      1956.0      0.1238      0.1866
## 84300903      152.50      1709.0      0.1444      0.4245
## 84348301      98.87      567.7      0.2098      0.8663
## 84358402      152.20      1575.0      0.1374      0.2050
## 843786      103.40      741.6      0.1791      0.5249
##      concavity_worst concave.points_worst symmetry_worst
## 842302      0.7119      0.2654      0.4601
## 842517      0.2416      0.1860      0.2750
## 84300903      0.4504      0.2430      0.3613
## 84348301      0.6869      0.2575      0.6638
## 84358402      0.4000      0.1625      0.2364
## 843786      0.5355      0.1741      0.3985
##      fractal_dimension_worst
## 842302      0.11890
## 842517      0.08902
## 84300903      0.08758
## 84348301      0.17300
## 84358402      0.07678
## 843786      0.12440
```

```
# also create a factor using only the diagnosis column
diagnosis <- factor(wisc.df[,1])
head(diagnosis)
```

```
## [1] M M M M M M
## Levels: B M
```

```
dim(wisc.data)
```

Q1. How many observations are in this dataset?

```
## [1] 569 30
```

The data frame has **569** observations.

```
sum(diagnosis == "M")
```

Q2. How many of the observations have a malignant diagnosis?

```
## [1] 212
```

Out of the 569 observations, **212** are malignant.

```
# save column names as a new vector  
wisc.colnames <- c(colnames(wisc.df))  
wisc.colnames
```

Q3. How many variables/features in the data are suffixed with `_mean`?

```
## [1] "diagnosis"           "radius_mean"  
## [3] "texture_mean"        "perimeter_mean"  
## [5] "area_mean"           "smoothness_mean"  
## [7] "compactness_mean"    "concavity_mean"  
## [9] "concave.points_mean" "symmetry_mean"  
## [11] "fractal_dimension_mean" "radius_se"  
## [13] "texture_se"          "perimeter_se"  
## [15] "area_se"             "smoothness_se"  
## [17] "compactness_se"      "concavity_se"  
## [19] "concave.points_se"   "symmetry_se"  
## [21] "fractal_dimension_se" "radius_worst"  
## [23] "texture_worst"       "perimeter_worst"  
## [25] "area_worst"          "smoothness_worst"  
## [27] "compactness_worst"   "concavity_worst"  
## [29] "concave.points_worst" "symmetry_worst"  
## [31] "fractal_dimension_worst"
```

```
# find the number of elements in wisc.colnames that contains _mean  
length(grep("_mean", wisc.colnames))
```

```
## [1] 10
```

10 features from the data are suffixed with `"_mean"`.

2. Principle Component Analysis

```
# Check column means and standard deviations
colMeans(wisc.data)
```

```
##          radius_mean      texture_mean      perimeter_mean
##      1.412729e+01      1.928965e+01      9.196903e+01
##          area_mean      smoothness_mean      compactness_mean
##      6.548891e+02      9.636028e-02      1.043410e-01
##      concavity_mean      concave.points_mean      symmetry_mean
##      8.879932e-02      4.891915e-02      1.811619e-01
## fractal_dimension_mean      radius_se      texture_se
##      6.279761e-02      4.051721e-01      1.216853e+00
##      perimeter_se      area_se      smoothness_se
##      2.866059e+00      4.033708e+01      7.040979e-03
##      compactness_se      concavity_se      concave.points_se
##      2.547814e-02      3.189372e-02      1.179614e-02
##      symmetry_se      fractal_dimension_se      radius_worst
##      2.054230e-02      3.794904e-03      1.626919e+01
##      texture_worst      perimeter_worst      area_worst
##      2.567722e+01      1.072612e+02      8.805831e+02
##      smoothness_worst      compactness_worst      concavity_worst
##      1.323686e-01      2.542650e-01      2.721885e-01
##      concave.points_worst      symmetry_worst      fractal_dimension_worst
##      1.146062e-01      2.900756e-01      8.394582e-02
```

```
apply(wisc.data,2,sd)
```

```
##          radius_mean      texture_mean      perimeter_mean
##      3.524049e+00      4.301036e+00      2.429898e+01
##          area_mean      smoothness_mean      compactness_mean
##      3.519141e+02      1.406413e-02      5.281276e-02
##      concavity_mean      concave.points_mean      symmetry_mean
##      7.971981e-02      3.880284e-02      2.741428e-02
## fractal_dimension_mean      radius_se      texture_se
##      7.060363e-03      2.773127e-01      5.516484e-01
##      perimeter_se      area_se      smoothness_se
##      2.021855e+00      4.549101e+01      3.002518e-03
##      compactness_se      concavity_se      concave.points_se
##      1.790818e-02      3.018606e-02      6.170285e-03
##      symmetry_se      fractal_dimension_se      radius_worst
##      8.266372e-03      2.646071e-03      4.833242e+00
##      texture_worst      perimeter_worst      area_worst
##      6.146258e+00      3.360254e+01      5.693570e+02
##      smoothness_worst      compactness_worst      concavity_worst
##      2.283243e-02      1.573365e-01      2.086243e-01
##      concave.points_worst      symmetry_worst      fractal_dimension_worst
##      6.573234e-02      6.186747e-02      1.806127e-02
```

```
# Perform PCA on wisc.data, since the data are on different magnitudes upon first inspection, we set scale = TRUE
wisc.pr <- prcomp( wisc.data, center = TRUE, scale. = TRUE )
summary(wisc.pr)
```

```
## Importance of components:
```

```

##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##          PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation  0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##          PC15     PC16     PC17     PC18     PC19     PC20     PC21
## Standard deviation  0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##          PC22     PC23     PC24     PC25     PC26     PC27     PC28
## Standard deviation  0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##          PC29     PC30
## Standard deviation  0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000


```

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)? 44.27%

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data? 3 (cumulative proportion achieves 72.6% at PC3)


Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data? 7 (cumulative proportion achieves 91% at PC7)

```
biplot(wisc.pr)
```




Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why? The plot is very difficult to understand because everything is on top of each other and we can't really see what is going on. (30 PCs are too many!)

```
# Scatter plot observations by components 1 and 2
plot( wisc.pr$x[,1], wisc.pr$x[,2] , col = diagnosis ,
      xlab = "PC1", ylab = "PC2")
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots? We see an overall separation of “M”(red) from “B”(black) samples. The plot of PC1 vs. PC2 has better separation on the y-axis than that of the PC1 vs. PC3 plot. This is reasonable because PC2 captures more variance in the data than PC3.

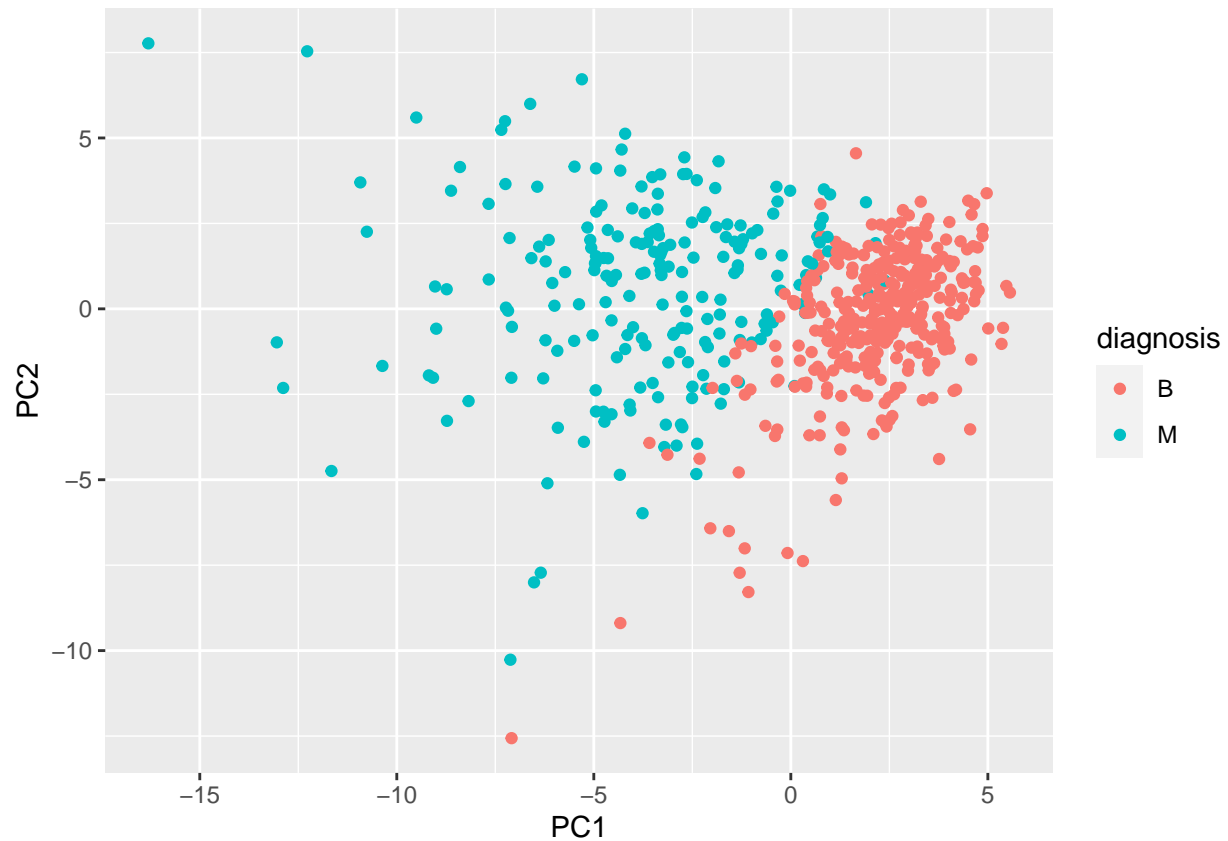
```
# Repeat for components 1 and 3
plot(wisc.pr$x[,1], wisc.pr$x[,3], col = diagnosis,
     xlab = "PC1", ylab = "PC3")
```

```
# Create a data.frame for ggplot
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

# Load the ggplot2 package
library(ggplot2)

# Make a scatter plot colored by diagnosis
ggplot(df) +
  aes(PC1, PC2, col= diagnosis) +
  geom_point()
```




```
# Calculate variance of each component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

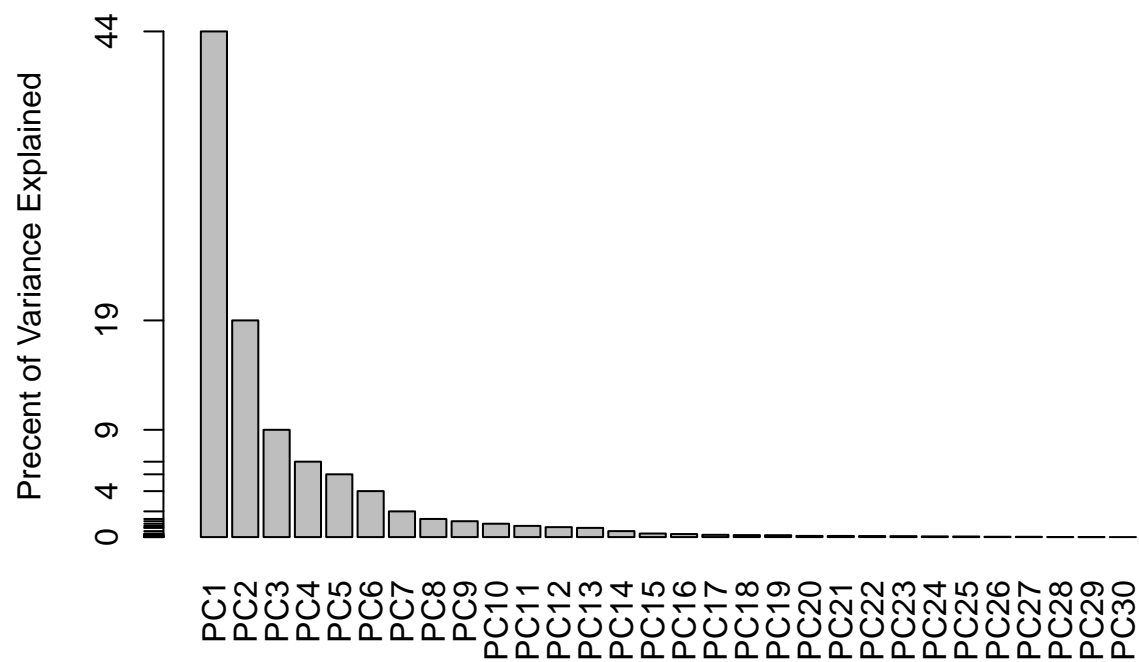
```
## [1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

```
# Variance explained by each principal component: pve
pve <- pr.var / sum(pr.var)

# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



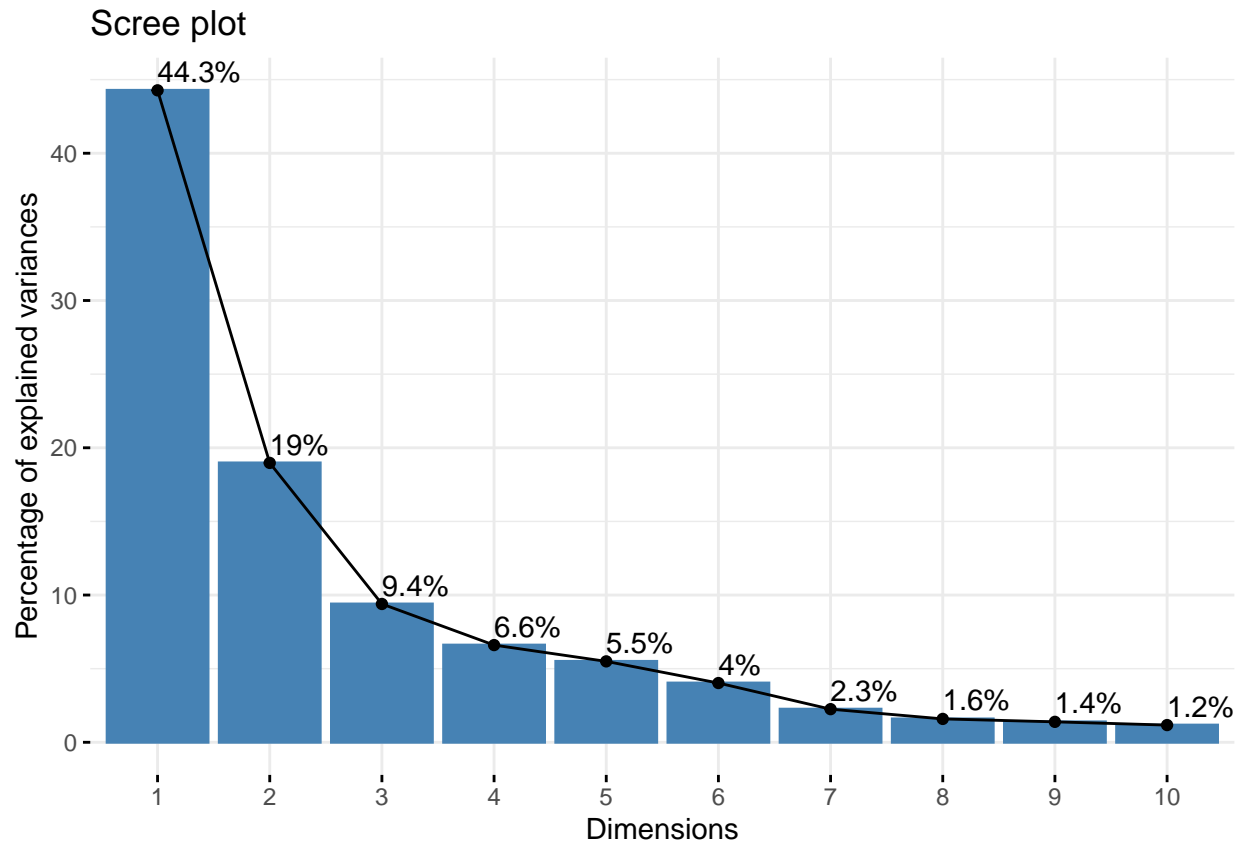
```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



```
## ggplot based graph
#install.packages("factoextra")
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
fviz_eig(wisc.pr, addlabels = TRUE)
```



```
wisc.pr$rotation["concave.points_mean",1]
```

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
## [1] -0.2608538
```

The loading vector of `concave.points_mean` for PC1 is about **-0.26**.

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data? Based on the scree plots and the variance table, **5 PCs** are required to explain 80% of the data.

3. Hierarchical Clustering


```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)
```

```
# Calculate the (Euclidean) distances between all
# pairs of observations in the new scaled dataset
# and assign the result to data.dist
data.dist <- dist(data.scaled)
```

```
# Create a hierarchical clustering model using
# complete linkage. Manually specify the method
# argument to hclust() and assign the results to
# wisc.hclust
wisc.hclust.complete <- hclust(data.dist, method = "complete")
```

Q11. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters? As shown, at height 19, the model has 4 clusters.

```
plot(wisc.hclust.complete)
abline(h = 15, col="blue", lty=2)
abline(h = 17, col="green", lty=2)
abline(h = 19, col="red", lty=2)
abline(h = 20, col="purple", lty=2)
```



```
data.dist
hclust (*, "complete")
```

```
wisc.hclust.clusters.complete <- cutree(wisc.hclust.complete, k = 4)
table(wisc.hclust.clusters.complete, diagnosis)
```

```
##                               diagnosis
```

```
## wisc.hclust.clusters.complete   B   M
##                               1  12 165
##                               2   2   5
##                               3 343  40
##                               4   0   2
```

```
wisc.hclust.clusters.complete9 <- cutree(wisc.hclust.complete, k = 9)
table(wisc.hclust.clusters.complete9, diagnosis)
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?


```
##                               diagnosis
## wisc.hclust.clusters.complete9   B   M
##                               1  12  86
##                               2   0  79
##                               3   0   3
##                               4 331  39
##                               5   2   0
##                               6  12   0
##                               7   0   2
##                               8   0   2
##                               9   0   1
```

Using $k = 9$, the result is slightly better.

Q13. Which method gives your favorite results for the same `data.dist` dataset? Explain your reasoning. Side-note: The method="ward.D2" creates groups such that variance is minimized within clusters. This has the effect of looking for spherical clusters with the process starting with all points in individual clusters (bottom up) and then repeatedly merging a pair of clusters such that when merged there is a minimum increase in total within-cluster variance. This process continues until a single group including all points (the top of the tree) is defined.

```
wisc.hclust <- hclust(data.dist, method = "centroid")
plot(wisc.hclust)
```

Cluster Dendrogram




data.dist
hclust (*, "centroid")

```
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)
table(wisc.hclust.clusters, diagnosis)
```

```
##           diagnosis
## wisc.hclust.clusters  B  M
##           1 357 208
##           2   0   1
##           3   0   1
##           4   0   2
```

```
wisc.hclust <- hclust(data.dist, method = "median")
plot(wisc.hclust)
```


Cluster Dendrogram




```
data.dist
hclust (*, "median")
```

```
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1 355 210
##              2   1   0
##              3   0   2
##              4   1   0
```

```
wisc.hclust <- hclust(data.dist, method = "ward.D2")
plot(wisc.hclust)
```

Cluster Dendrogram



```
data.dist
hclust (*, "ward.D2")
```

```
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##                   1   0 115
##                   2   6  48
##                   3 337  48
##                   4  14   1
```

I like the “**ward.D2**” method the best for it gives the neatest graph. The bottom-up grouping is also suitable for this dataset for we are eventually looking at either benign/malignant samples. However, the “complete” method seems to give the cleanest separation of the two sample types.

4. K-means Clustering

```
# Create a k-means model on wisc.data, assigning the result to wisc.km.
wisc.km <- kmeans(scale(wisc.data), centers= 2, nstart= 20)

# Use the table() function to compare the cluster membership of the k-means model
table(wisc.km$cluster, diagnosis)
```

```
##      diagnosis
```

```
##      B   M
##    1 14 175
##    2 343  37
```

```
table(wisc.hclust.clusters.complete, wisc.km$cluster)
```

```
##
## wisc.hclust.clusters.complete    1    2
##                                1 160  17
##                                2   7   0
##                                3  20 363
##                                4   2   0
```

Q14. How well does k-means separate the two diagnoses? How does it compare to your hclust results? k-means seem to have done a better job than hclust by giving two clusters with very good separation based on diagnosis.


5. Combining Methods

Using the minimum number of principal components required to describe at least 90% of the variability in the data, create a hierarchical clustering model with the linkage method="ward.D2". We use Ward's criterion here because it is based on multidimensional variance like principal components analysis. Assign the results to wisc.pr.hclust.

```
# Using the minimum number of principal components
# required to describe at least 90% of the
# variability in the data (PC1-7), create a
# hierarchical clustering model with the linkage
# method="ward.D2". We use Ward's criterion here
# because it is based on multidimensional variance
# like principal components analysis. Assign the
# results to wisc.pr.hclust.

wisc.pr.hclust <- hclust(dist(wisc.pr$x[, 1:7]), method = "ward.D2")
plot(wisc.pr.hclust)
```

Cluster Dendrogram



```
dist(wisc.pr$x[, 1:7])
hclust (*, "ward.D2")
```


```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
## grps
## 1 2
## 216 353
```


```
table(grps, diagnosis)
```

```
##      diagnosis
## grps    B    M
## 1    28 188
## 2   329  24
```

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```




```
# Note the color swap here as the hclust cluster 1
# is mostly "M" and cluster 2 is mostly "B" as we
# saw from the results of calling table(grps,
# diagnosis). To match things up we can turn our
# groups into a factor and reorder the levels so
# cluster 2 comes first and thus gets the first
# color (black) and cluster 1 gets the second color
# (red).
g <- as.factor(grps)
levels(g)
```

```
## [1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```


```
## [1] "2" "1"
```

```
# Plot using our re-ordered factor
plot(wisc.pr$x[,1:2], col=g)
```



```
# Fancy 3D plot
library(rgl)
plot3d(wisc.pr$x[,1:3], xlab="PC 1", ylab="PC 2", zlab="PC 3", cex=1.5, size=1, type="s", col=grps)
rglwidget(width = 400, height = 400)
```

```
## Warning in snapshot3d(scene = x, width = width, height = height): webshot = TRUE
## requires the webshot2 package; using rgl.snapshot() instead
```



```
# Use the distance along the first 7 PCs for
# clustering i.e. wisc.pr$x[, 1:7]
wisc.pr.hclust <- hclust(dist(wisc.pr$x[, 1:7]), method="ward.D2")

# Cut this hierarchical clustering model into 2
# clusters and assign the results to
# wisc.pr.hclust.clusters.
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)

# Compare to actual diagnoses
table(wisc.pr.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.pr.hclust.clusters  B  M
##              1  28 188
##              2 329  24
```

Q15. How well does the newly created model with two clusters separate out the two diagnoses?
 It did a better job than before and is pretty similar to the k-means results.


```
table(wisc.km$cluster, diagnosis)
```

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.km$cluster` and `wisc.hclust.clusters`) with the vector containing the actual diagnoses.

```
##      diagnosis
##      B      M
##  1  14 175
##  2 343  37
```

```
table(wisc.hclust.clusters.complete, diagnosis)
```

```
##                  diagnosis
## wisc.hclust.clusters.complete  B  M
##                  1 12 165
##                  2   2   5
##                  3 343  40
##                  4   0   2
```

Before PCA, the clustering is a lot messier and requires more than two clusters to achieve decent separation of the 2 diagnosis. After PCA, the separation is much cleaner.

6. Sensitivity/Specificity

Q17. Which of your analysis procedures resulted in a clustering model with the best specificity? How about sensitivity? **Sensitivity** refers to a test's ability to correctly detect ill patients who do have the condition. In our example here the sensitivity is the total number of samples in the cluster identified as predominantly malignant (cancerous) divided by the total number of known malignant samples. In other words: $TP/(TP+FN)$.

```
TPFN <- sum(diagnosis == "M")
kmeans.sensi = 175/TPFN
prWard2.sensi = 188/TPFN
```

Specificity relates to a test's ability to correctly reject healthy patients without a condition. In our example specificity is the proportion of benign (not cancerous) samples in the cluster identified as predominantly benign that are known to be benign. In other words: $TN/(TN+FN)$.

```
TNFN <- sum(diagnosis == "B")
kmeans.speci = 343/TNFN
prWard2.speci = 329/TNFN
```

Based on the above calculations, the kmeans clustering has higher specificity while the h-clust after PCA method has higher sensitivity.

7. Prediction


```

#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc

##          PC1          PC2          PC3          PC4          PC5          PC6          PC7
## [1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
## [2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
##          PC8          PC9          PC10          PC11          PC12          PC13          PC14
## [1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
## [2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
##          PC15          PC16          PC17          PC18          PC19          PC20
## [1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
## [2,] 0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
##          PC21          PC22          PC23          PC24          PC25          PC26
## [1,] 0.1228233 0.09358453 0.08347651  0.1223396  0.02124121 0.078884581
## [2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
##          PC27          PC28          PC29          PC30
## [1,] 0.220199544 -0.02946023 -0.015620933  0.005269029
## [2,] -0.001134152 0.09638361  0.002795349 -0.019015820

plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")

```



Q18. Which of these new patients should we prioritize for follow up based on your results?
Based on our clustering, patent 2 is located closer to the center of the malignant cluster, therefore, we should give priority to **patient 2**.